

▼ XGBoosting

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 %matplotlib inline
6 from sklearn.model_selection import GridSearchCV
7 from __future__ import print_function
8 import os
9
10 from google.colab import drive
11 drive.mount('/content/drive')
12
13 filepath = '/content/drive/My Drive/Colab Notebooks/Dataset/Walmart.csv'
14 xg = pd.read_csv(filepath)

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

```
1 xg.head()
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
0	1	05-02-2010	1643690.90	0	42.31	2.572	211.096358	8.106
1	1	12-02-2010	1641957.44	1	38.51	2.548	211.242170	8.106
2	1	19-02-2010	1611968.17	0	39.93	2.514	211.289143	8.106
3	1	26-02-2010	1409727.59	0	46.63	2.561	211.319643	8.106
4	1	05-03-2010	1554806.68	0	46.50	2.625	211.350143	8.106

```
1 xg.shape
```

(6435, 8)

```
1 xg.dtypes.value_counts()
```

float64	5
int64	2
object	1
dtype:	int64

```
1 xg = xg.drop('Date',axis=1)
```

```

1 float_columns = (xg.dtypes == np.float)
2 print( (xg.loc[:,float_columns].max()==1.0).all() )
3 print( (xg.loc[:,float_columns].min()==-1.0).all() )

```

```

False
False
<ipython-input-7-c6b27d4c13c4>:2: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use the built-in type name `float` instead.
Deprecation in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
float_columns = (xg.dtypes == np.float)

```

```

1 from sklearn.preprocessing import LabelEncoder
2
3 le = LabelEncoder()
4
5 xg['Holiday_Flag'] = le.fit_transform(xg['Holiday_Flag'])
6
7 le.classes_

```

array([0, 1])

```
1 xg.Holiday_Flag.unique()
```

array([0, 1])

```

1 from sklearn.model_selection import train_test_split
2
3 feature_columns = [x for x in xg.columns if x != 'Holiday_Flag']
4

```


```
5 X_train, X_test, y_train, y_test = train_test_split(xg[feature_columns], xg['Holiday_Flag'],  
6 test_size=0.3, random_state=42)
```

```
((4504, 6), (4504,), (1931, 6), (1931,))
```

```
1 from sklearn.ensemble import GradientBoostingClassifier
2 from sklearn.metrics import accuracy_score
3
4 error_list = list()
5
6 tree_list = [15, 50, 100, 200, 400]
7 for n_trees in tree_list:
8
9     GBC = GradientBoostingClassifier(n_estimators=n_trees,
10                                         subsample=0.5,
11                                         max_features=4,
12                                         random_state=42)
13
14     GBC.fit(X_train.values, y_train.values)
15     y_pred = GBC.predict(X_test)
16
17     error = 1. - accuracy_score(y_test, y_pred)
18
19     error_list.append(pd.Series({'n_trees': n_trees, 'error': error}))
20
21 error_xg = pd.concat(error_list, axis=1).T.set_index('n_trees')
22
23 error_xg
```

n_trees	
15.0	0.073019
50.0	0.069912
100.0	0.069394
200.0	0.064215
400.0	0.050751

```
1 sns.set_context('talk')
2 sns.set_style('white')
3 sns.set_palette('dark')
4
5 ax = error_xg.plot(marker='o')
6
7 ax.set(xlabel='n_trees', ylabel='error')
8 ax.set_xlim(0, max(error_xg.index)*1.1);
```



Number of trees vs error graph

```

1 param_grid = {'n_estimators': [200, 400],
2                 'learning_rate': [0.1, 0.01]}
3
4 GV_GBC = GridSearchCV(GradientBoostingClassifier(subsample=0.5,
5                                     max_features=4,
6                                     random_state=42),
7                                     param_grid=param_grid,
8                                     scoring='accuracy',
9                                     n_jobs=-1)
10
11 GV_GBC = GV_GBC.fit(X_train, y_train)

```

```

1 GV_GBC.best_estimator_
GradientBoostingClassifier(max_features=4, n_estimators=400, random_state=42,
                           subsample=0.5)

```

Accuracy, Classification Report, ROC AUC Score, Confusion Matrix for XGBoosting

```

1 from sklearn.metrics import classification_report
2
3 y_pred = GV_GBC.predict(X_test)
4 print(classification_report(y_pred, y_test))

precision    recall  f1-score   support

          0       1.00      0.95      0.97     1880
          1       0.33      0.92      0.49       51

   accuracy                           0.95      1931
  macro avg       0.67      0.94      0.73      1931
weighted avg       0.98      0.95      0.96      1931


```

- Based from the classification report, the team got an accuracy of 0.95 and with 1.00, 0.95 and 0.97 for precision, recall and f1-score, respectively.

```

1 from sklearn.metrics import confusion_matrix
2
3 sns.set_context('talk')
4 cm = confusion_matrix(y_test, y_pred)
5 ax = sns.heatmap(cm, annot=True, fmt='d')

```



```

1 from scipy.stats import uniform, randint
2 import xgboost as xgb
3 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, roc_auc_score

```

```

1 xgb_model = xgb.XGBClassifier(objective="binary:logistic", voting='hard', random_state=42)
2 xgb_model.fit(X_train, y_train)
3
4 y_pred = xgb_model.predict(X_test)
5
6 print(confusion_matrix(y_test, y_pred))

```

```

[[1789    1]
 [ 130   11]]

```

```

1 y_pred = xgb_model.predict(X_test)
2 print(classification_report(y_pred, y_test))

```

	precision	recall	f1-score	support
0	1.00	0.93	0.96	1919
1	0.08	0.92	0.14	12
accuracy			0.93	1931
macro avg	0.54	0.92	0.55	1931
weighted avg	0.99	0.93	0.96	1931

```
1 print(accuracy_score(y_test, y_pred))
```

```
0.9321595028482651
```

```
1 print(roc_auc_score(y_pred,y_test))
```

```
0.9244615250998783
```


After importing the xgboost as xgb and applying the code above, the classification report has slightly changed as well as as the accuracy. It has reduced to 0.93 now which is still a good value and the team got 0.924 as the roc auc score as well.

Agglomerative Clustering

```
1 from sklearn.cluster import KMeans
2
3 float_columns = [x for x in xg.columns if x not in ['Holiday_Flag', 'Unemployment']]
4
5 km = KMeans(n_clusters=2, random_state=42)
6 km = km.fit(xg[float_columns])
7
8 xg['kmeans'] = km.predict(xg[float_columns])
```

```
1 km_list = list()
2
3 for clust in range(1,21):
4     km = KMeans(n_clusters=clust, random_state=42)
5     km = km.fit(xg[float_columns])
6
7     km_list.append(pd.Series({'clusters': clust,
8                               'inertia': km.inertia_,
9                               'model': km}))
```

```
1 plot_xg = (pd.concat(km_list, axis=1)
2             .T
3             [['clusters','inertia']]
4             .set_index('clusters'))
5
6 ax = plot_xg.plot(marker='o',ls='--')
7 ax.set_xticks(range(0,21,2))
8 ax.set_xlim(0,21)
9 ax.set_xlabel('Cluster', ylabel='Inertia');
```



Inertia vs. Cluster graph


```
1 from sklearn.cluster import AgglomerativeClustering
2 from scipy.cluster import hierarchy
3 from matplotlib import colors
4
5 ag = AgglomerativeClustering(n_clusters=2, linkage='ward', compute_full_tree=True)
```

```
6 ag = ag.fit(xg[float_columns])
7 xg['agglom'] = ag.fit_predict(xg[float_columns])
```

```
1 (xg[['Holiday_Flag','agglom','kmeans']]  
2 .groupby(['Holiday_Flag','agglom','kmeans'])  
3 .size()  
4 .to_frame()  
5 .rename(columns={0:'number'}))
```

Holiday_Flag	agglom	kmeans	number
0	0	0	1384
		1	2375
1	1	0	2226
	0	0	95
1		1	197
	1	0	158

- ▼ Dendrogram produced by agglomerative clustering



```
1 from sklearn.ensemble import RandomForestClassifier
2 from sklearn.metrics import classification_report, roc_auc_score
3 from sklearn.model_selection import StratifiedShuffleSplit
4
5
6 y = (xg['Unemployment'] > 7).astype(int)
7 X_with_kmeans = xg.drop(['agglom', 'Holiday_Flag', 'Unemployment'], axis=1)
8 X_without_kmeans = X_with_kmeans.drop('kmeans', axis=1)
9 sss = StratifiedShuffleSplit(n_splits=10, random_state=6532)
10
11
12 def get_avg_roc_10splits(estimator, X, y):
13     roc_auc_list = []
14     for train_index, test_index in sss.split(X, y):
15         X_train, X_test = X.iloc[train_index], X.iloc[test_index]
16         y_train, y_test = y.iloc[train_index], y.iloc[test_index]
```

```

17     estimator.fit(X_train, y_train)
18     y_predicted = estimator.predict(X_test)
19     y_scored = estimator.predict_proba(X_test)[:, 1]
20     roc_auc_list.append(roc_auc_score(y_test, y_scored))
21     return np.mean(roc_auc_list)
22 # return classification_report(y_test, y_predicted)
23
24
25 estimator = RandomForestClassifier()
26 roc_with_kmeans = get_avg_roc_10splits(estimator, X_with_kmeans, y)
27 roc_without_kmeans = get_avg_roc_10splits(estimator, X_without_kmeans, y)
28 print("Without kmeans cluster as input to Random Forest, roc-auc is \"{}\"".format(roc_without_kmeans))
29 print("Using kmeans cluster as input to Random Forest, roc-auc is \"{}\"".format(roc_with_kmeans))
30

```

Without kmeans cluster as input to Random Forest, roc-auc is "0.9994285576505761"
Using kmeans cluster as input to Random Forest, roc-auc is "0.9992308434199739"

▼ Pipeline

```

1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 %matplotlib inline
6
7 from google.colab import drive
8 drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

```

1 filepath = '/content/drive/My Drive/Colab Notebooks/Dataset/Walmart.csv'
2 xg = pd.read_csv(filepath)

```

1 xg

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
0	1	05-02-2010	1643690.90	0	42.31	2.572	211.096358	8.106
1	1	12-02-2010	1641957.44	1	38.51	2.548	211.242170	8.106
2	1	19-02-2010	1611968.17	0	39.93	2.514	211.289143	8.106
3	1	26-02-2010	1409727.59	0	46.63	2.561	211.319643	8.106
4	1	05-03-2010	1554806.68	0	46.50	2.625	211.350143	8.106
...
6430	45	28-09-2012	713173.95	0	64.88	3.997	192.013558	8.684
6431	45	05-10-2012	733455.07	0	64.89	3.985	192.170412	8.667
6432	45	12-10-2012	734464.36	0	54.47	4.000	192.327265	8.667
6433	45	19-10-2012	718125.53	0	56.47	3.969	192.330854	8.667
6434	45	26-10-2012	760281.43	0	58.85	3.882	192.308899	8.667

6435 rows × 8 columns

```
1 xg = xg.drop('Date', axis=1)
```

```

1 for col in xg.columns:
2     xg[col] = xg[col].astype(np.float)

```

<ipython-input-56-6dbce3b12eef>:2: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence
Deprecation in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>
xg[col] = xg[col].astype(np.float)

```
1 xg_orig = xg.copy()
```

```

1 corr_mat = xg.corr()
2
3 # Strip the diagonal for future examination
4 for x in range(corr_mat.shape[0]):
5     corr_mat.iloc[x,x] = 0.0

```

```
6
7 corr_mat
```

	Store	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
Store	0.000000e+00	-0.335332	-4.386841e-16	-0.022659	0.060023	-0.209492	0.223531
Weekly_Sales	-3.353320e-01	0.000000	3.689097e-02	-0.063810	0.009464	-0.072634	-0.106176
Holiday_Flag	-4.386841e-16	0.036891	0.000000e+00	-0.155091	-0.078347	-0.002162	0.010960
Temperature	-2.265908e-02	-0.063810	-1.550913e-01	0.000000	0.144982	0.176888	0.101158
Fuel_Price	6.002295e-02	0.009464	-7.834652e-02	0.144982	0.000000	-0.170642	-0.034684
CPI	-2.094919e-01	-0.072634	-2.162091e-03	0.176888	-0.170642	0.000000	-0.302020
Unemployment	2.235313e-01	-0.106176	1.096028e-02	0.101158	-0.034684	-0.302020	0.000000

```
1 corr_mat.abs().idxmax()
```

```
Store      Weekly_Sales
Weekly_Sales      Store
Holiday_Flag      Temperature
Temperature      CPI
Fuel_Price        CPI
CPI              Unemployment
Unemployment     CPI
dtype: object
```

```
1 log_columns = xg.skew().sort_values(ascending=False)
2 log_columns = log_columns.loc[log_columns > 0.75]
3
4 log_columns
```

```
Holiday_Flag    3.373499
Unemployment   1.188144
dtype: float64
```

```
1 for col in log_columns.index:
2     xg[col] = np.log1p(xg[col])
```

```
1 from sklearn.preprocessing import MinMaxScaler
2
3 mms = MinMaxScaler()
4
5 for col in xg.columns:
6     xg[col] = mms.fit_transform(xg[[col]]).squeeze()
```

```
1 from sklearn.preprocessing import FunctionTransformer
2 from sklearn.pipeline import Pipeline
3
4 # The custom NumPy log transformer
5 log_transformer = FunctionTransformer(np.log1p)
6
7 # The pipeline
8 estimators = [('log1p', log_transformer), ('minmaxscale', MinMaxScaler())]
9 pipeline = Pipeline(estimators)
10
11 # Convert the original data
12 xg_pipe = pipeline.fit_transform(xg_orig)
```

```
1 np.allclose(xg_pipe, xg)
```

```
False
```

```
1 from sklearn.decomposition import PCA
2
3 pca_list = list()
4 feature_weight_list = list()
5
6 # Fit a range of PCA models
7
8 for n in range(1, 6):
9
10     # Create and fit the model
11     PCAmad = PCA(n_components=n)
12     PCAmad.fit(xg)
13
14     # Store the model and variance
15     pca_list.append(pd.Series({'n':n, 'model':PCAmad,
```

```

16             'var': PCAmod.explained_variance_ratio_.sum()})))
17
18     # Calculate and store feature importances
19     abs_feature_values = np.abs(PCAmod.components_).sum(axis=0)
20     feature_weight_list.append(pd.DataFrame({'n':n,
21                                              'features': xg.columns,
22                                              'values':abs_feature_values/abs_feature_values.sum()})))
23
24 pca_df = pd.concat(pca_list, axis=1).T.set_index('n')
25 pca_df

```

	model	var
n		
1	PCA(n_components=1)	0.37273
2	PCA(n_components=2)	0.563031
3	PCA(n_components=3)	0.717861
4	PCA(n_components=4)	0.832671
5	PCA(n_components=5)	0.90656

```

1 features_df = (pd.concat(feature_weight_list)
2                  .pivot(index='n', columns='features', values='values'))
3
4 features_df


```

	features	CPI	Fuel_Price	Holiday_Flag	Store	Temperature	Unemployment	Weekly_Sales
n								
1	0.577673	0.080874	0.001512	0.196542	0.050184	0.088784	0.004432	
2	0.372387	0.047145	0.012918	0.364854	0.045827	0.075438	0.081430	
3	0.252680	0.101460	0.193040	0.257552	0.079229	0.057383	0.058657	
4	0.201596	0.209181	0.201394	0.193538	0.085942	0.063331	0.045017	
5	0.160937	0.175849	0.173813	0.168016	0.153904	0.123667	0.043813	

```

1 sns.set_context('talk')
2
3 ax = pca_df['var'].plot(kind='bar')
4
5 ax.set(xlabel='Number of dimensions',
6        ylabel='Percent explained variance',
7        title='Explained Variance vs Dimensions');

```




This graph is the variance vs dimensions. As seen above, it is increasing consistently.

```

1 ax = features_df.plot(kind='bar')
2
3 ax.set(xlabel='Number of dimensions',
4        ylabel='Relative importance',
5        title='Feature importance vs Dimensions');

```



This graph shows the features as well as the number of dimensions vs the importance of the features with the features having different colors for the legends of the graph.


```

1 from sklearn.pipeline import Pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.model_selection import StratifiedShuffleSplit
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.metrics import accuracy_score
6 from sklearn.preprocessing import MinMaxScaler
7
8
9 X = xg.drop('Holiday_Flag', axis=1)
10 y = xg.Holiday_Flag
11 sss = StratifiedShuffleSplit(n_splits=10, random_state=42)
12
13 def get_avg_score(n):
14     pipe = [
15         ('scaler', MinMaxScaler()),
16         ('pca', PCA(n_components=n)),
17         ('estimator', LogisticRegression())
18     ]
19     pipe = Pipeline(pipe)
20     scores = []
21     for train_index, test_index in sss.split(X, y):
22         X_train, X_test = X.loc[train_index], X.loc[test_index]
23         y_train, y_test = y.loc[train_index], y.loc[test_index]
24         pipe.fit(X_train, y_train)
25         scores.append(accuracy_score(y_test, pipe.predict(X_test)))
26     return np.mean(scores)
27
28
29 ns = [1,2,3,4,5]
30 score_list = [get_avg_score(n) for n in ns]
```

```

1 sns.set_context('talk')
2
3 ax = plt.axes()
4 ax.plot(ns, score_list)
5 ax.set(xlabel='Number of Dimensions',
6        ylabel='Average Accuracy',
7        title='LogisticRegression Accuracy vs Number of dimensions on the Human Activity Dataset')
8 ax.grid(True)
```

LogisticRegression Accuracy vs Number of dimensions on the Human Activity Dataset



1 Start coding or generate with AI.

The graph shows the average accuracy of our data vs the number of dimensions.

```
1 xg.head()
```

	Store	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
0	0.0	0.397291	0.0	0.434149	0.050100	0.840500	0.545562
1	0.0	0.396811	1.0	0.396967	0.038076	0.841941	0.545562
2	0.0	0.388501	0.0	0.410861	0.021042	0.842405	0.545562
3	0.0	0.332458	0.0	0.476419	0.044589	0.842707	0.545562
4	0.0	0.372661	0.0	0.475147	0.076653	0.843008	0.545562